

# CS747 Assignment 1

Yash Salunkhe - 210020156

## 1 Task 1

My implementations of the algorithms UCB, KL-UCB and Thompson Sampling are in the file `task1.py`.

### 1.1 UCB

The UCB(Upper Confidence Bounds) algorithm achieves sub-linear regret with the regret being approximately 1400 after 250000 turns. It is performing better than the  $\epsilon$ -greedy algorithms in terms of regret achieved( $\epsilon$ -greedy gives linear regret).




Figure 1: Regret vs Horizon plot for the UCB algorithm

**Implementation** - Initialized 3 arrays and time variable in the `__init__` function.

`counts[i]` = Number of times the  $i^{\text{th}}$  arm has been pulled

`values[i]` = Empirical mean of the  $i^{\text{th}}$  arm

`UCB_bounds[i]` = The UCB value of the  $i^{\text{th}}$  arm

`time` = Total number of arms pulled

The `give_pull` function returns the index of the arm having the highest UCB bound.

The `get_reward` function updates the counts of the arm pulled, time and UCB bound for every arm. I have added  $1e-6$  to the counts(in the denominator) in the UCB bound to not encounter any division by 0 errors.

## 1.2 KL-UCB

The KL-UCB(Upper Confidence Bounds) algorithm achieves sub-linear regret with the maximum regret being approximately 110 after 250000 turns. It is performing better than the  $\epsilon$ -greedy algorithms in terms of regret achieved and has a tighter regret bound than the UCB algorithm.




Figure 2: Regret vs Horizon plot for the KL-UCB algorithm

**Implementation** - Initialized 3 arrays and time variable in the `__init__` function.

`counts[i]` = Number of times the  $i^{\text{th}}$  arm has been pulled

`values[i]` = Empirical mean of the  $i^{\text{th}}$  arm

`bounds[i]` = The KL-UCB value of the  $i^{\text{th}}$  arm ( $c = 0$ )

`time` = Total number of arms pulled

The `give_pull` function returns the index of the arm having the highest KL-UCB bound.

The `get_reward` function updates the counts of the arm pulled, time and KL-UCB bound for every arm. A binary search is applied to find the optimal  $q$  such that the KL equation holds.

For KL-divergence, I have used constants 0.0001 and 1.0000001 to handle edge cases like  $x = 0$  and 1.

### 1.3 Thompson Sampling

The Thompson Sampling algorithm achieves sub-linear regret with the maximum regret being approximately 140 after 250000 turns. It is performing better than the  $\epsilon$ -greedy algorithms in terms of regret achieved and has a tighter regret bound than the UCB, KL-UCB algorithm.



Figure 3: Regret vs Horizon plot for the Thompson Sampling algorithm

**Implementation** - Initialized 2 arrays in the `__init__` function.

`alpha[i]` = Number of successes(reward = 1) for the  $i^{\text{th}}$  arm.

`beta[i]` = Number of failures(reward = 0) for the  $i^{\text{th}}$  arm.

These have been initialized by 1 so that the sampling from the beta distribution does not give any errors.

The `give_pull` function returns the index of the arm having the highest Thompson Sampling value.

The `get_reward` function updates the `alpha`(if reward = 1) and `beta`(if reward = 0) for the arm.

## 2 Task 2

### 2.1 Task 2a

The plot obtained for regret vs  $p_2$  variation is as shown. The curve attains a maximum at a point between 0.8 and 0.9 and falls to zero at 0.9.

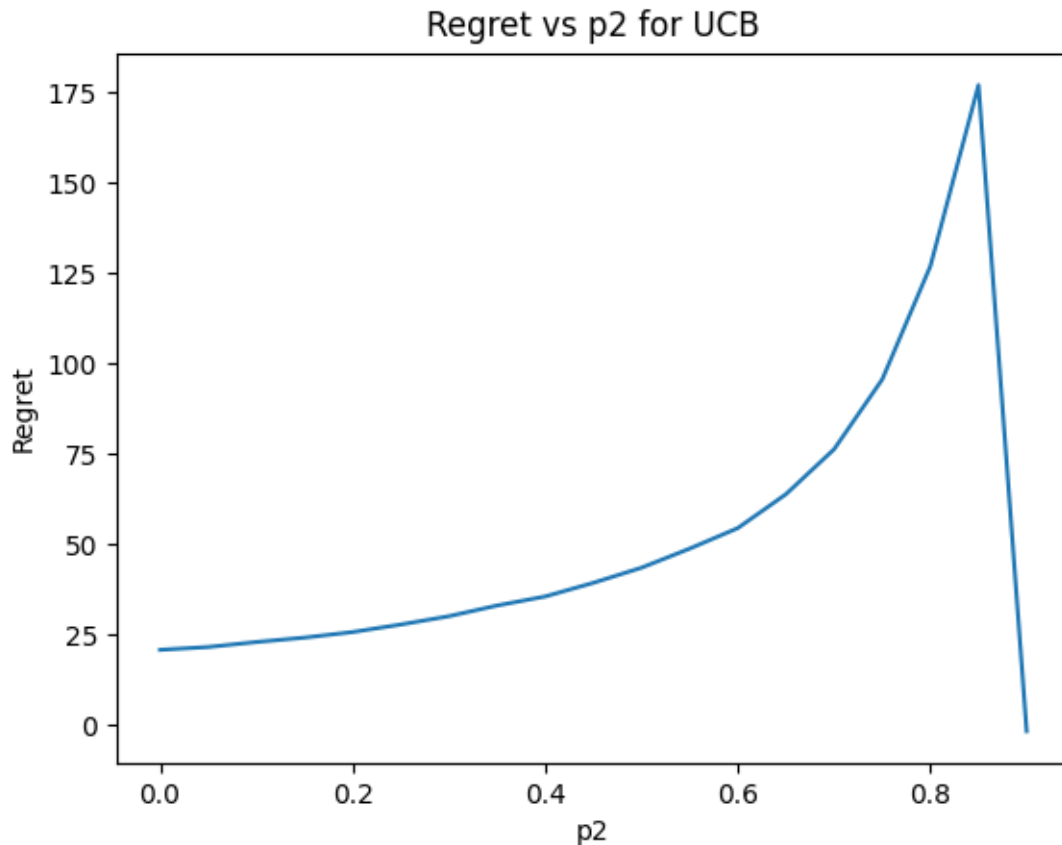


Figure 4: Variation of  $p_2$  with regret in UCB( $p_1 = 0.9$ )

**Reason for this trend** - As we increase  $p_2$  from 0 to 0.85, the time required for recognizing the optimal arm (the arm with probability  $p_1$ ) will increase. Sampling the arms with probabilities  $[0.9, 0]$  would take lesser time than  $[0.9, 0.8]$  to recognize that the first arm is optimal. Due to this, there would be a considerable number of pulls of the second arm which would result in increase in regret. When  $p_2$  reaches 0.9, every arm pulled is optimal, hence the expected regret goes to 0.

## 2.2 Task 2b

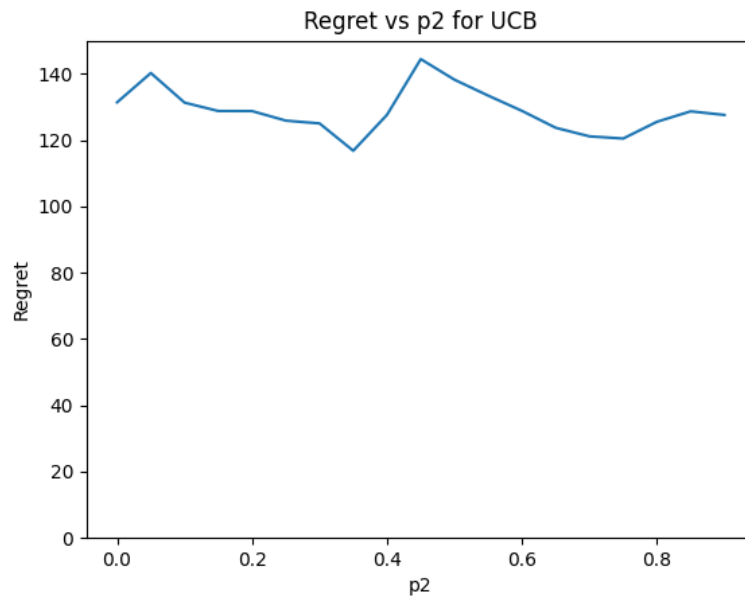


Figure 5: Variation of p2 with regret in UCB( $p_1 = p_2 + 0.1$ )




Figure 6: Variation of p2 with regret in KL-UCB( $p_1 = p_2 + 0.1$ )

**Reason for trends -**

In the UCB algorithm, a big difference is not seen in the regret as we change the value of  $p_2$ . The theoretical explanation for this is the expected regret for UCB is  $O(\frac{\log(T)}{p^* - p})$ . Since  $p^* - p$  is constant(=0.1), the expected regret does not fluctuate much.

For KL-UCB, there is a significant fluctuation in the regret as we tweak the value of  $p_2$  signifying that the regret is proportional to  $c \log(T)$  where  $c$  is the constant from Lai and Robbins' bound.

$$\frac{Regret}{\log(T)} \leq \sum_{i: p_i(l) \neq p^*(l)} \frac{p^*(l) - p(l)}{KL(p(l), p^*(l))},$$

Here, the regret does not solely depend on  $p^* - p$ , and hence, it changes significantly with change in  $p_2$ .

### 3 Task 3

Here, I used Thompson Sampling in a faulty bandit setting as my main goal is to maximize my reward. I have no control over my faulty bandit settings, I can maximize the output I get out of the non-faulty pulls which is achieved by Thompson Sampling. I preferred Thompson Sampling over UCB, KL-UCB as it has a tighter regret bound as seen in task 1.

The file `task3.py` contains the entire implementation.




Figure 7: Variation in reward vs horizon for my algorithm

This plot was obtained by adding `plt.plot(horizons,rewards)` to the `task3` function in `simulator.py`

## 4 Task 4

The file `task4.py` contains the entire implementation. Here, I again used Thompson Sampling in a multi-bandit setting as my main goal is to maximize my reward. I have no control over which bandit is being chosen as it is sampled randomly. Performing Thompson Sampling over the chosen bandit will maximize my reward. I preferred Thompson Sampling over UCB, KL-UCB as it has a tighter regret bound as seen in task 1.




Figure 8: Variation in reward vs horizon for my algorithm

This plot was obtained by adding `plt.plot(horizons,rewards)` to the `task4` function in `simulator.py`